DOCWEAVE / EXECUTED CORE PIPELINE

PDF / SCAN

Validate size and pages

Native text; OCR if empty

→

EXTRACT

Typed entity candidates

Rules / logistic / BiLSTM

→

LINK + STORE

Subject-predicate-object

SQLite + page provenance

→

RETRIEVE + VIEW

TF-IDF + typed graph path

Source-linked interface

Evidence contract: original PDF hash + page + chunk ID + exact source quote on each relation

Optional, not live-evaluated: Ollama LLM extraction | Neo4j export | Docling Markdown conversion